



# DM & ML

# Basic Methods

**Gergely Lukács**

Pázmány Péter Catholic University

Faculty of Information Technology  
and Bionics

Budapest, Hungary

lukacs@itk.ppke.hu

# Simplicity first

- Different kinds of simple structures/algorithms exist:
  - One attribute
  - All attributes independently
  - A linear combination
  - An instance-based representation
- Simple algorithms often work surprisingly well
- Useful for understanding data
- Basis for more complicated algorithms

1R

(One Rule:  
decision tree with  
only one level)


# 1R: Output format

- 1R: „one rule”
  - for classification
  - set of „rules” that all test one particular attribute  
(== 1-level „decision tree” – see later; „decision stump” )

X Y (class)

Outlook	Temp.	Humidity	Windy	Play
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

IF outlook=sunny THEN play=yes  
IF outlook=overcast THEN play=yes  
IF outlook=rainy THEN play=no



# 1R: Algorithm

- Basic version (assuming nominal attributes)
  - Create a ruleset for each of the attributes
    - One branch for each of the attribute's values
    - Each branch assigns most frequent class
    - Error rate: proportion of instances that don't belong to the majority class of their corresponding branch
  - Choose attribute and the corresponding ruleset with lowest error rate

# 1R Example: weather attributes

Outlook	Temp.	Humidity	Windy	Play
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

Attribute	Rules	Errors	Total errors
Outlook	Sunny → No	2/5	4/14
	Overcast → Yes	0/4	
	Rainy → Yes	2/5	
Temperature	Hot → No*	2/4	5/14
	Mild → Yes	2/6	
	Cool → Yes	1/4	
Humidity	High → No	3/7	4/14
	Normal → Yes	1/7	
Windy	False → Yes	2/8	5/14
	True → No*	3/6	

***\* A random choice is made between 2 equal outcomes***

# Pseudo-code for 1R

For each attribute,

    For each value of the attribute, make a rule as follows:

        count how often each class appears

        find the most frequent class

        make the rule assign that class to this attribute-value

    Calculate the error rate of the rules

Choose the rules with the smallest error rate

- “missing” is treated as a separate attribute value

# Dealing with numeric attributes

- Numeric attributes are discretized: the range of the attribute is divided into a set of intervals:
  - Instances are sorted according to attribute's values
  - Breakpoints are placed where the class changes (so that the total error is minimized)
    - Equal values, different classes: majority class
- Example: *temperature* from weather data

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	Yes	No	No	/Yes	Yes	Yes	No	Yes	Yes	No



# The problem of overfitting

- Discretization procedure is very sensitive to noise
  - A single instance with an incorrect class label will most likely result in a separate interval
- Simple solution: **enforce minimum number of instances in majority class per interval**
- Weather data example (with minimum set to 3):

64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	/No	/Yes	Yes	Yes	No	No	/Yes	Yes	Yes	No	/Yes	Yes	/No
64	65	68	69	70	71	72	72	75	75	80	81	83	85
Yes	No	Yes	Yes	Yes	No	No	Yes	Yes	Yes	No	Yes	Yes	No

# Result of overfitting avoidance


- Resulting rule sets:

Attribute	Rules	Errors	Total errors
Outlook	Sunny $\rightarrow$ No	2/5	4/14
	Overcast $\rightarrow$ Yes	0/4	
	Rainy $\rightarrow$ Yes	2/5	
Temperature	$\leq 77.5 \rightarrow$ Yes	3/10	5/14
	$> 77.5 \rightarrow$ No*	2/4	
Humidity	$\leq 82.5 \rightarrow$ Yes	1/7	3/14
	$> 82.5$ and $\leq 95.5 \rightarrow$ No	2/6	
	$> 95.5 \rightarrow$ Yes	0/1	
Windy	False $\rightarrow$ Yes	2/8	5/14
	True $\rightarrow$ No*	3/6	

***\* A random choice is made between 2 equal outcomes***


# Discussion of 1R

- 1R was described in a paper by Holte (1993)
  - Contains an experimental evaluation on 16 datasets (using *cross-validation* so that results were representative of performance on future data)
  - After some experimentation the minimum number of instances in majority class was set to 6
- Simplicity first ...
- Data investigation ...  
how important individual attributes are
- Part of more complex algorithms  
(ensemble learning)
- Decision boundary
  - illustration with 2 input variables (two axes),  
binary classification (class: colour)



# Naïve Bayes

# NB: The model



- “Opposite” of 1R: use all attributes
- Assumption: Attributes are *statistically independent*, **given the class value**
  - This means that knowledge about the value of a particular attribute doesn’t tell us anything about the value of another attribute -- **if the class is known**


- 

- Fixed

$$P(x|y) = P(x_1, x_2, \dots | y) = P(x_1|y)P(x_2|y)\dots = \prod_{i=1}^n P(x_i|y)$$

- Although assumption almost never (entirely) correct, this scheme works quite often well in practice!

# Likelihood vs probability




- **StatQuest: Probability vs Likelihood**  
<https://www.youtube.com/watch?v=pYxNSUDSFH4>

# Maximum Likelihood Estimation

- **maximum likelihood estimation (MLE)** of parameters given observations. MLE attempts to find the parameter values that maximize the likelihood function, given the observations
- Naive Bayes: learning: estimating parameter values of the model given data, based on the MLE  
(Details: <http://www.datasciencecourse.org/notes/mle/>)

# Naive Bayes: learning (training)



- **(Relative) frequencies of values in training data**
  - For binary classifier:

$$P(y = k) = \frac{1}{m} \times \#\{j | y^{(j)} = k\}$$

and

$$P(x_i = l | y = k) = \frac{\#\{j | y^{(j)} = k \text{ and } x_i^{(j)} = l\}}{\#\{j | y^{(j)} = k\}}$$

$$k \in \{0, 1\}, l \in \llbracket 1, L \rrbracket$$



# Probabilities for the weather data: Frequencies of cases

$X_1$		$X_4$		$y$
Outlook	Temp.	Humidity	Windy	Play
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

Outlook			Temperature			Humidity			Windy		Play		
	Yes	No		Yes	No		Yes	No		Yes	No	Yes	No
Sunny	2	3	Hot	2	2	High	3	4	False	6	2	9	5
Overcast	4	0	Mild	4	2	Normal	6	1	True	3	3		
Rainy	3	2	Cool	3	1								

# Probabilities for the weather data: probability estimations

Outlook			Temperature			Humidity			Windy			Play	
	Yes	No		Yes	No		Yes	No		Yes	No	Yes	No
Sunny	2	3	Hot	2	2	High	3	4	False	6	2	9	5
Overcast	4	0	Mild	4	2	Normal	6	1	True	3	3		
Rainy	3	2	Cool	3	1								
Sunny	2/9	3/5	Hot	2/9	2/5	High	3/9	4/5	False	6/9	2/5	9/14	5/14
Overcast	4/9	0/5	Mild	4/9	2/5	Normal	6/9	1/5	True	3/9	3/5		
Rainy	3/9	2/5	Cool	3/9	1/5								

$P(\text{Outlook}=\text{Sunny} \mid \text{Play}=\text{Yes})$

(Outlook:  $X_1$ ,  $E$ )

$P(\text{Play}=\text{Yes})$

(Play:  $Y$ , class,  $H$ )

Can you draw the Naive Bayes network,  
including the (conditional and apriori) probabilities?

# NB: prediction (new/unknown instances)

- Classification learning: what's the probability of the class given the input attributes?

# Probabilities for the weather data: Prediction task

Outlook			Temperature			Humidity			Windy		Play		
	Yes	No		Yes	No		Yes	No		Yes	No	Yes	No
Sunny	2	3	Hot	2	2	High	3	4	False	6	2	9	5
Overcast	4	0	Mild	4	2	Normal	6	1	True	3	3		
Rainy	3	2	Cool	3	1								
Sunny	2/9	3/5	Hot	2/9	2/5	High	3/9	4/5	False	6/9	2/5	9/14	5/14
Overcast	4/9	0/5	Mild	4/9	2/5	Normal	6/9	1/5	True	3/9	3/5		
Rainy	3/9	2/5	Cool	3/9	1/5								

- A new day:

Outlook	Temp.	Humidity	Windy	Play
Sunny	Cool	High	True	?

# The weather data example

Outlook	Temp.	Humidity	Windy	Play
Sunny	Cool	High	True	?

← **Evidence**  
 **$E (== X)$**

$$\Pr[\text{yes} | E] = \Pr[\text{Outlook} = \text{Sunny} | \text{yes}] \times$$

$$\Pr[\text{Temperature} = \text{Cool} | \text{yes}] \times$$

$$\Pr[\text{Humidity} = \text{High} | \text{yes}] \times$$

$$\Pr[\text{Windy} = \text{True} | \text{yes}] \times \frac{\Pr[\text{yes}]}{\Pr[E]}$$

$$= \frac{2/9 \times 3/9 \times 3/9 \times 3/9 \times 9/14}{\Pr[E]}$$

**Cond. probability  
for class “yes”**

$P(\text{Play}=\text{yes}|\text{Outlook}=\text{Sunny}$   
 $\wedge \dots)$

# Probabilities for the weather data

Outlook			Temperature			Humidity			Windy		Play		
	Yes	No		Yes	No		Yes	No		Yes	No	Yes	No
Sunny	2	3	Hot	2	2	High	3	4	False	6	2	9	5
Overcast	4	0	Mild	4	2	Normal	6	1	True	3	3		
Rainy	3	2	Cool	3	1								
Sunny	2/9	3/5	Hot	2/9	2/5	High	3/9	4/5	False	6/9	2/5	9/14	5/14
Overcast	4/9	0/5	Mild	4/9	2/5	Normal	6/9	1/5	True	3/9	3/5		
Rainy	3/9	2/5	Cool	3/9	1/5								

- A new day:

Outlook	Temp.	Humidity	Windy	Play
Sunny	Cool	High	True	?

## Numerator

For “yes” =  $2/9 \times 3/9 \times 3/9 \times 3/9 \times 9/14 = 0.0053$

For “no” =  $3/5 \times 1/5 \times 4/5 \times 3/5 \times 5/14 = 0.0206$

Conversion into a probability by normalization (instead of by calculating the denominator,  $P(X)$ ):

$P(\text{“yes”}) = 0.0053 / (0.0053 + 0.0206) = 0.205$

$P(\text{“no”}) = 0.0206 / (0.0053 + 0.0206) = 0.795$

# The “zero-frequency problem”

- What if an attribute value doesn't occur with every class value (e.g. “Outlook = Overcast” for class “no”)?

- Probability will be zero!

$$\Pr[\text{Outlook} = \text{Overcast} \mid \text{no}] = 0$$

- *A posteriori* probability will also be zero!

$$\Pr[\text{no} \mid E] = 0$$

(No matter how likely the other values are!)

- Remedy: add 1 to the count of **every** value-class combination for such problematic attribute („Laplace estimator”)
- Result: probabilities will never be zero

# Applying Laplace estimator

Outlook			Temperature			Humidity			Windy			Play	
	Yes	No		Yes	No		Yes	No		Yes	No	Yes	No
Sunny	2	3	Hot	2	2	High	3	4	False	6	2	9	5
Overcast	4	0	Mild	4	2	Normal	6	1	True	3	3		
Rainy	3	2	Cool	3	1								
Sunny	2/9	3/5	Hot	2/9	2/5	High	3/9	4/5	False	6/9	2/5	9/14	5/14
Overcast	4/9	0/5	Mild	4/9	2/5	Normal	6/9	1/5	True	3/9	3/5		
Rainy	3/9	2/5	Cool	3/9	1/5								

Outlook			Temperature			Humidity			Windy			Play	
	Yes	No		Yes	No		Yes	No		Yes	No	Yes	No
Sunny	2	4	Hot	2	2	High	3	4	False	6	2	9	5
Overcast	4	1	Mild	4	2	Normal	6	1	True	3	3		
Rainy	3	3	Cool	3	1								
Sunny	2/9	4/8	Hot	2/9	2/5	High	3/9	4/5	False	6/9	2/5	9/14	5/14
Overcast	4/9	1/8	Mild	4/9	2/5	Normal	6/9	1/5	True	3/9	3/5		
Rainy	3/9	3/8	Cool	3/9	1/5								



# Missing values

- Training
  - instance is not included in frequency count for attribute-value/class combination
- Classification
  - attribute will be omitted from calculation

– Example:

Outlook	Temp.	Humidity	Windy	Play
?	Cool	High	True	?

Numerator for “yes” =  $3/9 \times 3/9 \times 3/9 \times 9/14 = 0.0238$

Numerator for “no” =  $1/5 \times 4/5 \times 3/5 \times 5/14 = 0.0343$

$P(\text{“yes”}) = 0.0238 / (0.0238 + 0.0343) = 41\%$

$P(\text{“no”}) = 0.0343 / (0.0238 + 0.0343) = 59\%$

# Dealing with numeric attributes

- Usual assumption: attributes have a *normal/Gaussian* probability distribution (given the class)
- The *probability density function* for the normal distribution is defined by two parameters:

- The *sample mean*  $\mu$ :

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i$$

- The *standard deviation*  $\sigma$ :

$$\sigma = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \mu)^2}$$

- The density function  $f(x)$ :

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

# Statistics for the weather data

Outlook			Temperature			Humidity		Windy		Play			
	Yes	No		Yes	No		Yes	No		Yes	No	Yes	No
Sunny	2	3		83	85		86	85	False	6	2	9	5
Overcast	4	0		70	80		96	90	True	3	3		
Rainy	3	2		68	65		80	70					
				64	72		65	95					
				69	71		70	91					
				....		....							
Sunny	2/9	3/5	mean	73	74.6	mean	79.1	86.2	False	6/9	2/5	9/14	5/14
Overcast	4/9	0/5	std dev	6.2	7.9	std dev	10.2	9.7	True	3/9	3/5		
Rainy	3/9	2/5											

- Example density value:

$$f(\text{temperature} = 66 | \text{yes}) = \frac{1}{\sqrt{2\pi} 6.2} e^{-\frac{(66-73)^2}{2 \cdot 6.2^2}} = 0.0340$$

# Probability densities

- Relationship between probability and density:

$$\Pr[a \leq x \leq b] = \int_a^b f(t) dt$$

$$\Pr[c - \frac{\varepsilon}{2} < x < c + \frac{\varepsilon}{2}] \approx \varepsilon \star f(c)$$

- But: this doesn't change calculation of *a posteriori* probabilities because  $\varepsilon$  cancels out

# Classifying a new day

- A new day:

Outlook	Temp.	Humidity	Windy	Play
Sunny	66	90	true	?

Numerator of “yes” =  $2/9 \times 0.0340 \times 0.0221 \times 3/9 \times 9/14 = 0.000036$

Numerator of “no” =  $3/5 \times 0.0291 \times 0.0380 \times 3/5 \times 5/14 = 0.000136$

$P(\text{“yes”}) = 0.000036 / (0.000036 + 0.000136) = 20.9\%$

$P(\text{“no”}) = 0.000136 / (0.000036 + 0.000136) = 79.1\%$


- Note also: many numeric attributes are not normally distributed
  - Known distribution: work with it
  - Unknown distribution: *kernel density estimators*
- Missing values during training: not included in calculation of mean and standard deviation

# Discussion of Naïve Bayes

- Naïve Bayes works surprisingly well (even if independence assumption is clearly violated)
- Too many **redundant attributes** (e.g. identical attributes) will cause problems
- **Irrelevant attributes** are OK, they are filtered out
- Fast to train, fast to classify
  - even for large number of attributes
  - incremental model building possible (streamed data)
- Outputs probability (?) distribution, not just a single class (+)
- (Does not learn training data perfectly (+/-))
- Typical application: text/spam detection

# Discussion of Naïve Bayes, 2

- Decision boundary....



# Divide-and-conquer

## Decision tree

### ID3

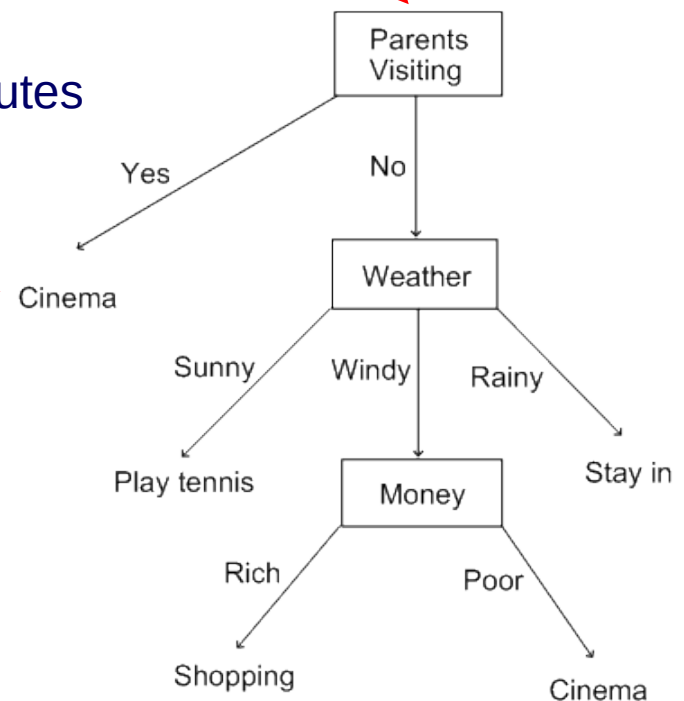


# **DECISION TREE – AS KNOWLEDGE REPRESENTATION**

# Decision trees

## (classification, regression)

- Nodes involve testing a particular attribute
  - usually, attribute value is compared to constant;
  - other possibilities:
    - comparing values of two attributes
    - using a function of one or more attributes
- Leaves assign
  - class value, or
  - set of class values, or
  - probability distribution to instances



# Testing nominal and numeric attributes

- Nominal attribute: number of children (branches) usually equal to number values  $\Rightarrow$  attribute won't get tested more than once
  - Other possibility: division into two subsets
- Numeric attribute: test whether value is greater or less than constant  $\Rightarrow$  attribute may get tested several times
  - Other possibility: three/multi-way split)
    - Integer: *less than, equal to, greater than*
    - Real: *below, within, above*

# Example: Labor negotiations data

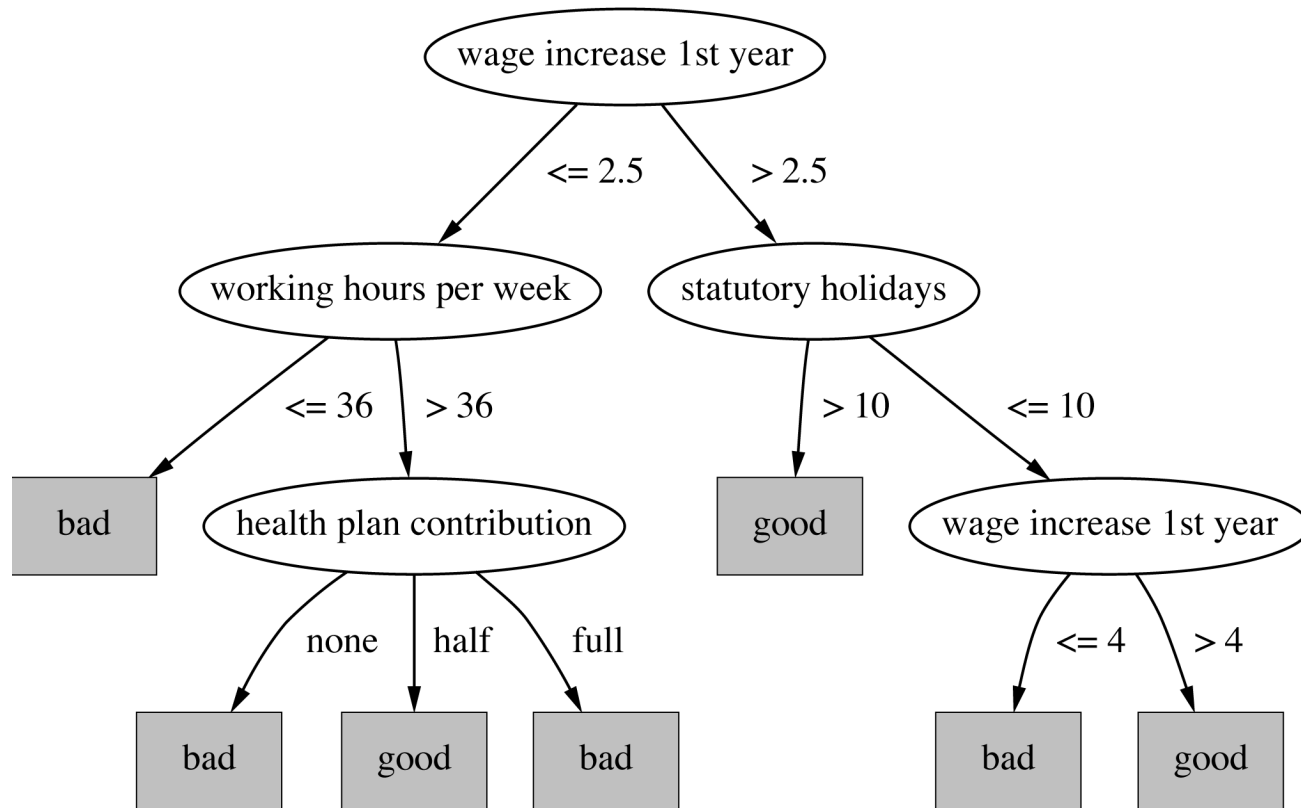
(Instances)

Attribute	Type	1	2	3	...	40
Duration	(Number of years)	1	2	3		2
Wage increase first year	Percentage	2%	4%	4.3%		4.5%
Wage increase second year	Percentage	?	5%	4.4%		4.0%
Wage increase third year	Percentage	?	?	?		?
Cost of living adjustment	{none,tcf,tc}	none	tcf	?		none
Working hours per week	(Number of hours)	28	35	38		40
Pension	{none,ret-allw, empl-cntr}	none	?	?		?
Standby pay	Percentage	?	13%	?		?
Shift-work supplement	Percentage	?	5%	4%		4%
Education allowance	{yes,no}	yes	?	?		?
Statutory holidays	(Number of days)	11	15	12		12
Vacation	{below-avg,avg,gen}	avg	gen	gen		avg
Long-term disability assistance	{yes,no}	no	?	?		yes
Dental plan contribution	{none,half,full}	none	?	full		full
Bereavement assistance	{yes,no}	no	?	?		yes
Health plan contribution	{none,half,full}	none	?	full		half
Acceptability of contract	{good,bad}	bad	good	good		good

(Attributes)

(Class attribute)

# Decision tree for labor negotiations data



# Missing values

- ***Does absence of value have some significance?***
- Yes  $\Rightarrow$  “missing” is a separate value
- No  $\Rightarrow$  “missing” must be treated in a special way
  - Solution A: assign instance to ***most popular branch***
  - Solution B: split instance into pieces
    - Pieces receive ***weight according to fraction of training instances*** that go down each branch

For both, we need some statistics on the training data in the tree  
– in addition to the previous tests

# **LEARNING A DECISION TREE**

## **– ID 3 ALGORITHM**

# Constructing decision trees

- Preferred: small decision tree (== small number of levels/attributes tested)
- Greedy algorithm
- Normal procedure: top down in recursive *divide-and-conquer* fashion
  - Attribute is selected for root node and branch is created for each possible attribute value
  - Instances are split into subsets (one for each branch extending from the node)
  - Procedure is repeated recursively for each branch, using only instances that reach the branch
- Process stops if all instances have the same class
  - or no further split makes them purer







# Example data

(class attr,  
y)

Outlook	Temp.	Humidity	Windy	Play
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

# Which attribute to select?



# A criterion for attribute selection

- Which is the best attribute?
  - The one which will result in the smallest tree
  - Heuristic: choose the attribute that produces the “purest” nodes
- Popular *impurity criterion*: **information gain**
  - Information gain increases with the average „purity” of the subsets that an attribute produces
- Strategy: choose attribute that results in greatest information gain
- „purity”: entropy

# Example: attribute *Outlook*

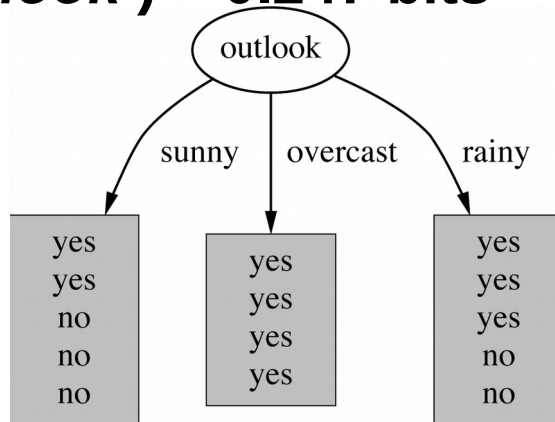
- *Outlook = Sunny* :  
$$\text{Info}([2,3]) = \text{entropy}(2/5, 3/5) =$$
$$= -2/5 \log_2(2/5) - 3/5 \log_2(3/5) = 0.971 \text{ bits}$$
- *Outlook = Overcast* :  
$$\text{Info}([4,0]) = \text{entropy}(1,0) = 0 \text{ bits}$$
- *Outlook = Rainy* :  
$$\text{Info}([3,2]) = \text{entropy}(3/5, 2/5) =$$
$$= -3/5 \log_2(3/5) - 2/5 \log_2(2/5) = 0.971 \text{ bits}$$
- Expected information for attribute  
(weighed sum, weighs: number of instances):  
$$\text{Info}([2,3], [4,0], [3,2]) = 5/14 * 0.971 + 4/14 * 0 + 5/14 * 0.971 =$$
$$= 0.693 \text{ bits}$$

# Computing information gain


- Information gain:  
information before splitting – information after splitting  
 $\text{gain}(\textit{Outlook}) = \text{info}([9,5]) - \text{info}([2,3],[4,0],[3,2])$   
 $= 0.940 - 0.693 = 0.247 \text{ bits}$

# Which attribute to select?


$\text{gain}(\text{Outlook}) = 0.247$  bits




$\text{gain}(\text{Humidity}) = 0.152$  bits



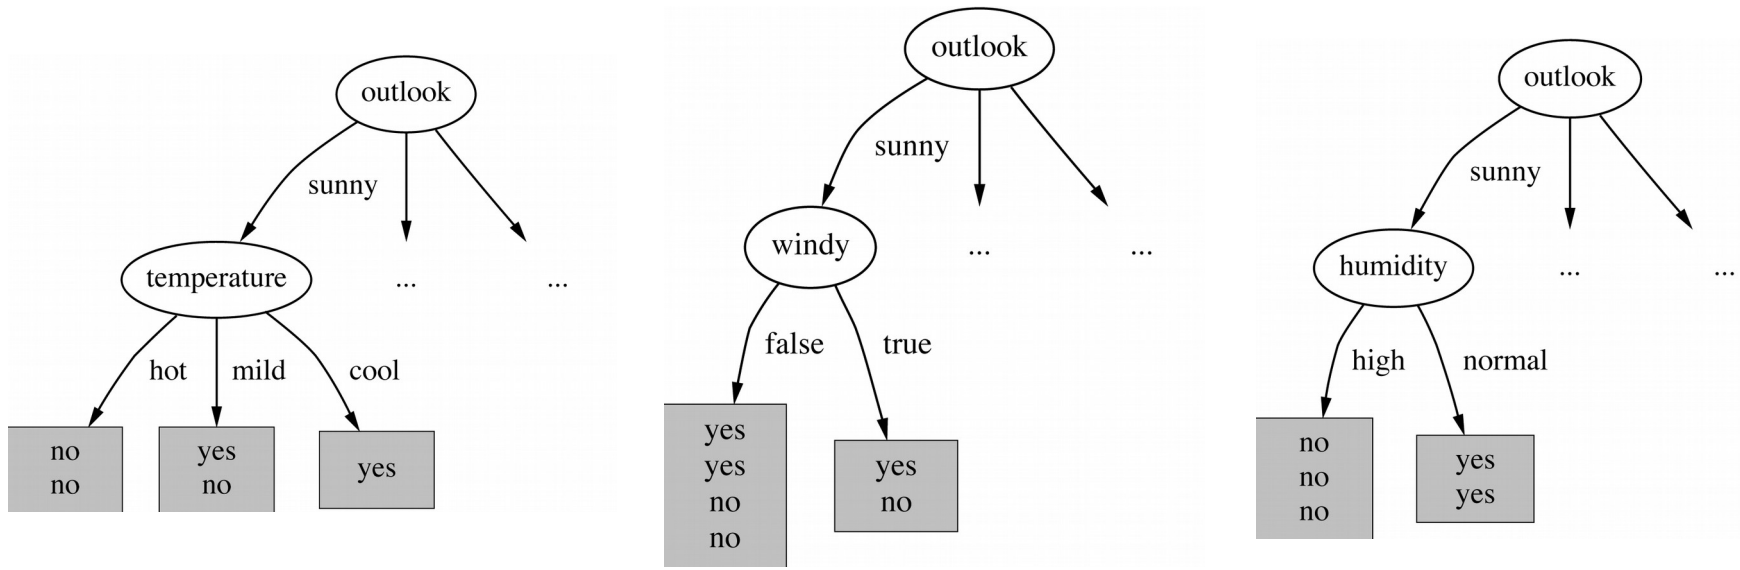
$\text{gain}(\text{Windy}) = 0.048$  bits



$\text{gain}(\text{Temperature}) = 0.029$  bits



# Continuing to split




$\text{gain}(\text{"Temperature"}) = 0.571 \text{ bits}$

$\text{gain}(\text{"Windy"}) = 0.019 \text{ bits}$

$\text{gain}(\text{"Humidity"}) = 0.971 \text{ bits}$

# The final decision tree



- Note: *not* all leaves need to be pure; sometimes identical instances have different classes  
⇒ Splitting stops when data can't be split any further



# Wish list for a purity measure

- Properties we require from a purity measure:
  - Pure node: zero
  - Impurity is maximal (i.e. all classes equally likely): maximal (e.g., binary case: 1)
  - ***Multistage property*** (i.e. decisions can be made in several stages):

$$\text{measure}([2,3,4]) = \text{measure}([2,7]) + (7/9) \times \text{measure}([3,4])$$

- **Entropy** is the only function that satisfies all three properties!

# Computing information

- Information is measured in *bits*
  - Given a probability distribution, the info required to predict an event is the distribution's *entropy*
  - Entropy gives the information required in bits (this can involve fractions of bits!)
- Formula for computing the entropy:

$$\text{entropy}(p_1, p_2, \dots, p_n) = -p_1 \log p_1 - p_2 \log p_2 - \dots - p_n \log p_n$$

# Some properties of the entropy

- The multistage property:

$$\text{entropy}(p, q, r) = \text{entropy}(p, q + r) + (q + r) \times \text{entropy}\left(\frac{q}{q + r}, \frac{r}{q + r}\right)$$

(where  $p + q + r = 1$ )

- Simplification of computation:

$$\begin{aligned} \text{info}([2, 3, 4]) &= -2/9 \times \log(2/9) - 3/9 \times \log(3/9) - 4/9 \times \log(4/9) \\ &= [-2\log 2 - 3\log 3 - 4\log 4 + 9\log 9]/9 \end{aligned}$$


# Highly-branching attributes

- Problematic: attributes with a large number of values (extreme case: ID code)
- Subsets are more likely to be pure if there is a large number of values
  - ⇒ Information gain is biased towards choosing attributes with a large number of values
  - ⇒ This may result in *overfitting* (selection of an attribute that is non-optimal for prediction)

# The weather data with ID code (extreme example)

ID code	Outlook	Temp.	Humidity	Windy	Play
A	Sunny	Hot	High	False	No
B	Sunny	Hot	High	True	No
C	Overcast	Hot	High	False	Yes
D	Rainy	Mild	High	False	Yes
E	Rainy	Cool	Normal	False	Yes
F	Rainy	Cool	Normal	True	No
G	Overcast	Cool	Normal	True	Yes
H	Sunny	Mild	High	False	No
I	Sunny	Cool	Normal	False	Yes
J	Rainy	Mild	Normal	False	Yes
K	Sunny	Mild	Normal	True	Yes
L	Overcast	Mild	High	True	Yes
M	Overcast	Hot	Normal	False	Yes
N	Rainy	Mild	High	True	No

# Tree stump for ID code attribute



- Entropy of split:

$\text{info}(\text{"ID code"}) = \text{info}([0,1]) + \text{info}([0,1]) + \dots + \text{info}([0,1]) = 0 \text{ bits}$

⇒ Information gain is maximal for ID code (namely 0.940 bits)

# The gain ratio

- *Gain ratio*: a modification of the information gain that reduces its bias towards highly branching attributes
- Gain ratio takes the number and the size of branches into account when choosing an attribute (disregarding the class values!)
  - It corrects the information gain by taking the *intrinsic information* of a split into account
- Intrinsic information: entropy of distributing instances into branches (i.e. how much info do we need to tell which branch an instance belongs to)

# Computing the gain ratio

- Example: intrinsic information for ID code

$$\text{info}([1,1,\dots,1]) = 14 \times (-1/14 \times \log 1/14) = 3.807 \text{ bits}$$

- Value of attribute decreases as intrinsic information gets larger
- Definition of gain ratio:

$$\text{gain\_ratio}(\text{"Attribute"}) = \frac{\text{gain}(\text{"Attribute"})}{\text{intrinsic\_info}(\text{"Attribute"})}$$

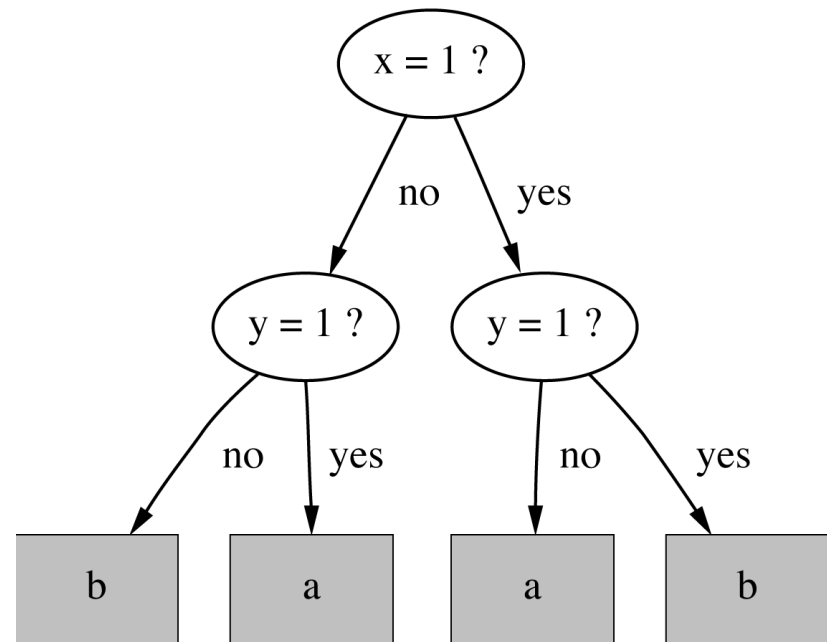

- Example:  $\text{gain\_ratio}(\text{"ID\_code"}) = \frac{0.940 \text{ bits}}{3.807 \text{ bits}} = 0.246$



# Gain ratios for weather data

Outlook		Temperature	
Info:	0.693	Info:	0.911
Gain: 0.940-0.693	0.247	Gain: 0.940-0.911	0.029
Split info: info([5,4,5])	1.577	Split info: info([4,6,4])	1.362
Gain ratio: 0.247/1.577	0.156	Gain ratio: 0.029/1.362	0.021
Humidity		Windy	
Info:	0.788	Info:	0.892
Gain: 0.940-0.788	0.152	Gain: 0.940-0.892	0.048
Split info: info([7,7])	1.000	Split info: info([8,6])	0.985
Gain ratio: 0.152/1	0.152	Gain ratio: 0.048/0.985	0.049

# Knowledge representation: decision tree (extension): XOR



- awkward representation
- hard to learn
- XOR problem fortunately rare

# Discussion

- Algorithm for top-down induction of decision trees (“ID3”) was developed by Ross Quinlan
  - Led to development of C4.5, which can deal with numeric attributes, missing values, and noisy data
- Similar approach: Regression tree (CART: Classification and Regression Tree)
- There are many other attribute selection criteria! (But almost no difference in accuracy of result.)
- Decision Tree: „white box”
- Decision Trees: traditionally most widely used ML models!


# Classification Rule learning

(Covering algorithm)  
(Generating rules)  
(PRISM)

# Knowledge Representation: Classification rules

- Popular alternative to decision trees
- *Antecedent* (pre-condition): a series of tests (just like the tests at the nodes of a decision tree)
  - tests are usually logically AND-ed together
  - but may also be general logical expressions
- *Consequent* (conclusion):
  - classes, or
  - set of classes, or
  - probability distribution assigned by rule
- Individual rules are often logically OR-ed together...

# Classification rule - example



# “Nuggets” of knowledge

- Are rules independent pieces of knowledge? (It seems easy to add a rule to an existing rule base.)
- Problem: ignores how rules are executed
- Two ways of executing a rule set:
  - Ordered set of rules (“decision list”)
    - Order is important for interpretation
  - Unordered set of rules
    - Rules may overlap and lead to different conclusions for the same instance

# Interpreting rules

- What if two or more rules conflict?
  - Give no conclusion at all?
  - Go with rule that is most popular on training data?
  - ...
- What if no rule applies to an instance?
  - Give no conclusion at all?
  - Go with class that is most frequent in training data?
  - ...



# Special case: boolean class

- Assumption: if instance does not belong to class “yes”, it belongs to class “no”
- Trick: only learn rules for class “yes” and use default rule for “no”

`If x = 1 and y = 1 then class = a`

`If z = 1 and w = 1 then class = a`

`Otherwise class = b`

- Order of rules is not important. No conflicts!

# From trees to rules

- Easy: converting a tree into a set of rules
  - one rule for each leaf:
    - antecedent contains a condition for every node on the path from the root to the leaf
    - consequent is class assigned by the leaf
- Produces rules that are unambiguous
  - doesn't matter in which order they are executed
- But: resulting rules are unnecessarily complex

# From rules to trees

- More difficult: transforming a rule set into a tree
  - tree cannot easily express disjunction between rules

# Linear models for regression (=predicting numeric value)

# Linear models

- Inputs (attribute values) and output are all numeric
- Output is the sum of weighted attribute values
  - $x = w_0 + w_1 a_1 + w_2 a_2 + \dots + w_k a_k$

The trick is to find good values for the weights


# Example:

## Predicting CPU performance

- 209 different computer configurations

	Cycle time (ns)	Main memory (Kb)		Cache (Kb)	Channels		Performance
	MYCT	MMIN	MMAX	CACH	CHMIN	CHMAX	PRP
1	125	256	6000	256	16	128	198
2	29	8000	32000	32	8	32	269
...							
208	480	512	8000	32	0	0	67
209	480	1000	4000	0	0	0	45

# A linear regression function for the CPU performance data



$$\begin{aligned} \text{PRP} = & \\ & - 56.1 \\ & + 0.049 \text{ MYCT} \\ & + 0.015 \text{ MMIN} \\ & + 0.006 \text{ MMAX} \\ & + 0.630 \text{ CACH} \\ & - 0.270 \text{ CHMIN} \\ & + 1.46 \text{ CHMAX} \end{aligned}$$

# Linear regression

- Outcome is linear combination of attributes

$$x = w_0 + w_1 a_1 + w_2 a_2 + \dots + w_k a_k$$

- Weights are calculated from the training data
- Predicted value for first training instance  $\mathbf{a}^{(1)}$

$$w_0 a_0^{(1)} + w_1 a_1^{(1)} + w_2 a_2^{(1)} + \dots + w_k a_k^{(1)} = \sum_{j=0}^k w_j a_j^{(1)}$$



# Goal: Minimum Mean Squared Error (MSE)

- *Mean Squared Error (MSE), or L2 loss*

$$\sum_{i=1}^n \left( x^{(i)} - \sum_{j=0}^k w_j a_j^{(i)} \right)^2$$

- $k+1$  coefficients are chosen so that MSE the training data is minimized
- Assumptions ...
  - Error: Normal distribution
- MSE sensitive to outliers!
- Minimization of *absolute error* is more difficult!


# Training: Minimizing Mean Squared Error (MSE)

- Coefficients can be derived using matrix inversion
- Iterative solution -- gradient decent -- computationally less expensive
- Can be done if there are more instances than attributes (roughly speaking)

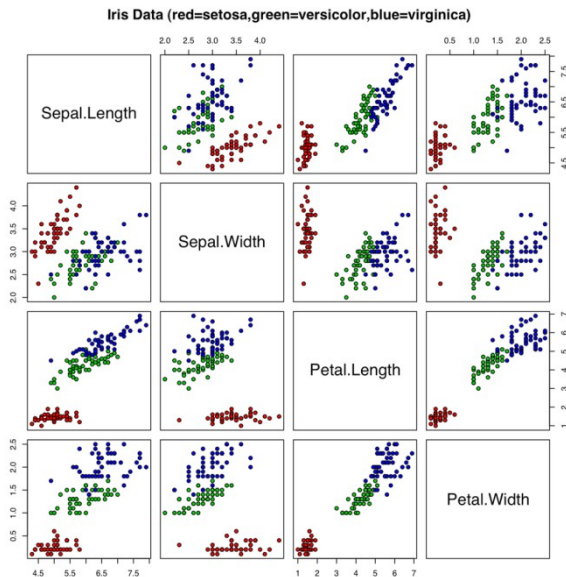
# Linear models for binary classification

- Binary classification
- Prediction is made by plugging in observed values of the attributes into the expression
  - Predict one class if output  $> 0$  (or 0.5, ...), and the other class if output  $< 0$  (or 0.5, ...)
- Decision boundary
  - defines where the decision changes from one class value to the other
  - hyperplane (high-dimensional plane)


# Example: Classifying iris flowers



	Sepal length	Sepal width	Petal length	Petal width	Type
1	5.1	3.5	1.4	0.2	Iris setosa
2	4.9	3.0	1.4	0.2	Iris setosa
...					
51	7.0	3.2	4.7	1.4	Iris versicolor
52	6.4	3.2	4.5	1.5	Iris versicolor
...					
	6.3	3.3	6.0	2.5	Iris virginica
	5.8	2.7	5.1	1.9	Iris virginica



# Separating setosas from versicolors



- **$2.0 - 0.5\text{PETAL-LENGTH} - 0.8\text{PETAL-WIDTH} = 0$**   
(2 input attributes)

# Classification

- Problems
  - membership values are not in  $[0,1]$  range, so aren't proper probability estimates
  - assumptions used in least-square regression
    - normal distribution with same standard deviation for 0 -1 values ??!!

# Linear models: logistic regression


- Linear model for a transformed target variable
- Assume we have two classes
- Logistic regression replaces the target

$$P[1|a_1, a_2, \dots, a_k]$$

- by this target

$$\log\left(\frac{P[1|a_1, a_2, \dots, a_k]}{(1 - P[1|a_1, a_2, \dots, a_k])}\right)$$

- Logit transformation maps  $[0,1]$  to  $(-\infty, +\infty)$   
(Inverse function of logit: sigmoid function)




# Logit transformation (Sigmoid function): Resulting model

$$Pr[1|a_1, a_2, \dots, a_k] = \frac{1}{(1 + e^{-w_0 - w_1 a_1 - \dots - w_k a_k})}$$

- Example

– Model with

- $w_0 = 0.5$  (moves curve left-right) and
- $w_1 = 1$  (steepness)






# Maximum likelihood

- Parameters are found from training data using *maximum likelihood*
- Aim: choosing parameters to maximize probability of training data („maximum likelihood”)
  - Independence: probabilities multiplied
  - Maximizing multiplied probabilities : maximizing sum of logarithms of probabilities

$$\sum_{i=1}^n (1 - x^{(i)}) \log(1 - \text{Pr}[1 | a_1^{(i)}, a_2^{(i)}, \dots, a_k^{(i)}]) + x^{(i)} \log \text{Pr}[1 | a_1^{(i)}, a_2^{(i)}, \dots, a_k^{(i)}]$$

where the  $x^{(i)}$  are either 0 or 1

Log loss  $\sum -y_i z_i + \log(1 + e^{z_i})$



# Logistic regression 2

- The independent variables should be independent of each other. That is, the model should have little or no multicollinearity.
- The independent variables are linearly related to logit function - probability.
- Logistic regression requires quite large sample sizes.

# Linear models are hyperplanes

- Decision boundary for two-class logistic regression is where probability equals 0.5:

$$Pr[1|a_1, a_2, \dots, a_k] = 1 / (1 + \exp(-w_0 - w_1 a_1 - \dots - w_k a_k)) = 0.5$$

- which occurs when  $-w_0 - w_1 a_1 - \dots - w_k a_k = 0$

# Linear models: the perceptron

- Don't actually need probability estimates if all we want to do is classification
- Different approach: learn separating hyperplane
- Assumption: data is *linearly separable*
- Hyperplane:  $0 = w_0 a_0 + w_1 a_1 + w_2 a_2 + \dots + w_k a_k$   
where we again assume that there is a constant attribute with value 1 (*bias*)
- If sum is greater than zero we predict the first class, otherwise the second class

# Perceptron learning rule

Set all weights to zero

Until all instances in the training data are classified correctly

For each instance  $I$  in the training data

If  $I$  is classified incorrectly by the perceptron

If  $I$  belongs to the first class add it to the weight vector

else subtract it from the weight vector

- Why does this work?

Consider situation where instance  $a$  pertaining to the first class has been added:

$$(w_0 + a_0)a_0 + (w_1 + a_1)a_1 + (w_2 + a_2)a_2 + \dots + (w_k + a_k)a_k$$

This means output for  $a$  has increased by:


$$a_0 a_0 + a_1 a_1 + a_2 a_2 + \dots + a_k a_k$$

This number is  $a_0 a_0 + a_1 a_1 + a_2 a_2 + \dots + a_k a_k$  has moved into the correct direction (and we can show output decreases for instances of other class)

- Convergence guaranteed -- if data linearly separable --, upper bound on number of steps

# Perceptron as a neural network

Output layer



Input layer

# Linear models: Winnow

- Binary classification
- Mistake-driven
- **Multiplicative updates** (perceptron: additive updates):
  - Weights are multiplied (or divided) by a user-specified parameter ( $\alpha > 1$ )
- User-specific threshold (perceptron: 0)
  - Predict first class, if:

$$w_0 a_0 + w_1 a_1 + w_2 a_2 + \dots + w_k a_k > \theta$$

# Winnow learning rule

```
while some instances are misclassified
  for each instance  $a$  in the training data
    classify  $a$  using the current weights
    if the predicted class is incorrect
      if  $a$  belongs to the first class
        for each  $\tilde{a}_i$  that is 1, multiply  $\tilde{w}_i$  by alpha
        (if  $\tilde{a}_i$  is 0, leave  $\tilde{w}_i$  unchanged)
      otherwise
        for each  $\tilde{a}_i$  that is 1, divide  $\tilde{w}_i$  by alpha
        (if  $\tilde{a}_i$  is 0, leave  $\tilde{w}_i$  unchanged)
```

- Effectively focuses on relevant features

Winnowing





# Instance-based learning

# Instance-based representation

- Simplest form of learning: *rote learning*
  - Training instances are searched for instance that most closely resembles new instance
  - The instances themselves represent the knowledge
    - ***Lazy learning!***
  - Also called *instance-based* learning
- Similarity function

# The distance function

- Simplest case: one numeric attribute
  - Distance is the difference between the two attribute values involved (or a function thereof)
- Several numeric attributes: typically Euclidean distance is used and attributes are **normalized**
- Taking  $\sqrt{(a_1^{(1)} - a_1^{(2)})^2 + (a_2^{(1)} - a_2^{(2)})^2 + \dots + (a_k^{(1)} - a_k^{(2)})^2}$  when comparing distances
- Other popular metric: *city-block metric*
  - Adds differences without squaring them

# The distance function

- Nominal attributes: distance is set to 1 if values are different, 0 if they are equal
- Are all attributes equally important?
  - weighting the attributes might be necessary

# Discussion of nearest-neighbor learning

- Statisticians have used  $k$ -NN since early 1950s
- Often very accurate


## Limitations (of the basic algorithm)

- White box, generalisation ?!
- Speed (of prediction!!)
  - E.g. 1E6 instances, 100 attributes (dimensions)
- Noisy data ?!
- Domains of attributes different ?!
  - comparing e.g.  $[0,1]$  with  $[-5000,5000]$  ?
- Irrelevant attributes (only noise, with no relevance)

# Clustering


# Clustering: Knowledge Representation

- Divide instances into “natural” groups (clusters)
  - Instances in a group (cluster) similar to each other, in some or the other way
  - No class attribute, unlabelled data
  - Unsupervised learning




# Representing clusters I

Simple 2-D representation



Venn diagram



**disjoint vs. overlapping**




# Representing clusters II

## Probabilistic assignment

	1	2	3
a	0.4	0.1	0.5
b	0.1	0.8	0.1
c	0.3	0.3	0.4
d	0.1	0.1	0.8
e	0.4	0.2	0.4
f	0.1	0.4	0.5
g	0.7	0.2	0.1
h	0.5	0.4	0.1
...			

**probabilistic**  
**(vs. deterministic)**

## Dendrogram



NB: dendron is the  
Greek word for tree

**hierarchical**  
**(vs. flat)**

# The $k$ -means algorithm

- disjoint, deterministic, and flat clusters
- number of clusters ( $k$ ) preselected

1. Choose  $k$  cluster centers (e.g. randomly)

2. Repeat until convergence

1. assign instances to clusters (based on distance; closest cluster center)
2. compute *centroids* of clusters

(convergence: assignment of instances to clusters do not change)


# K-means: illustration




(a)




(b)




(c)



(d)






(e)



(f)

# Discussion

- Algorithm minimizes squared distance to cluster centers
- Converges
  - distortion function (==sum of squared errors) converges
  - typically also cluster centers and assignment of instances to clusters)
- Result can vary significantly
  - based on initial choice of seeds  
(restarting with different random seeds)
- Can get trapped in local minimum  
(in practice not typical)
- Number of clusters?
- Distance?
- Speed of distance calculations...
- Probabilistic, hierarchical, overlapping...



# Mining Association Rules (Apriori Algorithm)

# Association rules (knowledge representation)

- Form of rules


If a and b and c occurs in the set, then t is also part of the set

- **if** Milk = true **and**  
    Bread = True  
**then** Eggs = true **and**  
    Coke = True

- Typically boolean attributes
  - Compact notation: **if** Milk and Bread **then** Eggs and Coke
- No selected class attribute
- Multiple attributes on consequent side possible/typical
- Rules interpreted individually!
  - **if** Milk **and** True **then** Eggs **and** Coke

# Association rule mining, applications

- Application areas
  - Market basket analysis
  - ~Sequence mining



<i>TID</i>	<i>Items</i>
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

$\{\text{Diaper}\} \rightarrow \{\text{Beer}\},$   
 $\{\text{Milk, Bread}\} \rightarrow \{\text{Eggs, Coke}\},$   
 $\{\text{Beer, Bread}\} \rightarrow \{\text{Milk}\},$

# Association rule mining

Immense number of potential rules...

Sufficiently **general** rules

Sufficiently **correct/precise** rules

a,b,c,d,e

b,c,d

a,b,c,d

a,b,c,d,e

a,c,e,f

d,c,e,f

a,b,c,d,f



**b,c,d** co-occur

in 71% of the cases



**If b and c then d**  
in 100% of the cases



# Support and confidence of a rule

- (Rule:  $X \rightarrow Y$ )
- *Support/coverage*: number of instances predicted correctly
  - $\text{Supp}(X \text{ and } Y)$  „Itemset“ – no matter, whether an item (attribute) is on the left or on the right hand side of the rule
  - Can also be calculated as a relative number (relative to the number of all instances)
- *Confidence/accuracy*: number of correct predictions, as proportion of all instances that rule applies to
  - $\text{confidence} = \text{Supp}(X \text{ and } Y) / \text{Supp}(X)$
- *Lift*: How popular is Y, when X, **considering also how popular Y is**
  - $\text{lift} = \text{Confidence of Rule} / \text{Relative support of Y}$
  - 1: no association,  
     $>1$  Y is more likely when X,  
     $<1$  Y is less likely when X


# Association rules

## Example for support & confidence

Outlook	Temp.	Humidity	Windy	Play
Sunny	Hot	High	False	No
Sunny	Hot	High	True	No
Overcast	Hot	High	False	Yes
Rainy	Mild	High	False	Yes
Rainy	Cool	Normal	False	Yes
Rainy	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Sunny	Mild	High	False	No
Sunny	Cool	Normal	False	Yes
Rainy	Mild	Normal	False	Yes
Sunny	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Rainy	Mild	High	True	No

If temperature = cool  
then humidity = normal

support = 4  
confidence = 100%



If temperature = cool  
then play = yes

support = 3  
confidence = 75%



# Interpreting association rules

Interpretation is not obvious:

```
If windy = false
    and play = no
then outlook = sunny
    and humidity = high
```

is *not* the same as the two rules

```
If windy = false
    and play = no
then outlook =
sunny
```

```
If windy = false
    and play = no
then humidity =
high
```

It means that the following **also** holds:

```
If humidity = high
    and windy = false
    and play = no
then outlook = sunny
```



Why?

# Interpreting association rules

Interpretation is not obvious:

```
If windy = false  
    and play = no  
then outlook = sunny  
    and humidity = high
```

is *not* the same as the two rules

```
If windy = false  
    and play = no  
then outlook =  
sunny
```

```
If windy = false  
    and play = no  
then humidity =  
high
```

It means that the following **also** holds:


```
If humidity = high  
    and windy = false  
    and play = no  
then outlook = sunny
```

Why?  
Support...  
Confidence...

# **ASSOCIATION RULES: APRIORI ALGORITHM**

# Mining association rules

- Normally: minimum support and confidence pre-specified (e.g. for weather data: support  $\geq 2$  and confidence  $\geq 95\%$ , -> 58 rules)
- Brute force: All rules, filtering on the bases of support and confidence
- Immense number of possible associations
- Computational complexity!



# Association rules

- But: we can look for high support rules (itemsets) and high confidence rules directly!
- Two steps
  1. High support (frequent item sets)
  2. High confidence

# Frequent Item sets

- Support: number of instances correctly covered by association rule
  - The same as the number of instances covered by *all* tests in the rule (LHS and RHS!)
- *Item*: one test (that is, attribute-value pair)
- *Item set*: all items occurring in a rule (LHS and RHS !!)
- Goal: only rules that exceed pre-defined support
  - ⇒ We can do it by finding all item sets with the given minimum support (=frequent item sets) and generating rules from them!



# Frequent Item Sets for weather data

One-item sets	Two-item sets	Three-item sets	Four-item sets
Outlook = Sunny (5)	Outlook = Sunny Temperature = Mild (2)	Outlook = Sunny Temperature = Hot Humidity = High (2)	Outlook = Sunny Temperature = Hot Humidity = High Play = No (2)
Temperature = Cool (4)	Outlook = Sunny Humidity = High (3)	Outlook = Sunny Humidity = High Windy = False (2)	Outlook = Rainy Temperature = Mild Windy = False Play = Yes (2)
...	...	...	...

- Number of frequent item sets (with minimum support of two) :
  - 12 with 1 item,
  - 47 with 2 items,
  - 39 with 3 items,
  - 6 with 4 items
  - 0 with 5 items

# Generating item sets efficiently

- How can we efficiently find all frequent item sets (item sets with support value higher than the threshold for support)?
- If  $(A \ B)$  is frequent item set, then  $(A)$  and  $(B)$  have to be frequent item sets as well!
  - Necessary, but not sufficient condition:  
if  $(A)$  and  $(B)$  are frequent item sets,  $(A \ B)$  is not necessarily frequent, too.
- In general: if  $X$  is frequent  $k$ -item set, then all  $(k-1)$ -item subsets of  $X$  are also frequent

# Generating item sets efficiently

- Find frequent (support larger than threshold) one-item sets
- Compute candidate frequent  $k$ -item sets by merging frequent  $(k-1)$ -item sets
- Check (real) support of candidate frequent  $k$ -items sets by reading through the whole database
- Remark: The algorithm reduces the number of  $k$ -item sets, the support of which has to be checked by reading through the whole database.

# An example

- Given: five frequent three-item sets

(A B C) , (A B D) , (A C D) , (A C E) , (B C D)

- Lexicographically ordered!

- Candidate frequent four-item sets:

(A B C D)      candidate

(B C D E)      no candidate, e.g., because  
of (C D E)

- Check real support by counting instances in dataset!

# Generating rules from an item set

- Once all item sets with minimum support have been generated, we can turn them into rules
- **Example:** Humidity = Normal, Windy = False, Play = Yes (4)
- Seven ( $2^N-1$ ) potential rules:

<b>If</b> Humidity = Normal <b>and</b> Windy = False	<b>then</b> Play = Yes	4/4
<b>If</b> Humidity = Normal <b>and</b> Play = Yes	<b>then</b> Windy = False	4/6
<b>If</b> Windy = False <b>and</b> Play = Yes	<b>then</b> Humidity = Normal	4/6
<b>If</b> Humidity = Normal	<b>then</b> Windy = False <b>and</b> Play = Yes	4/7
<b>If</b> Windy = False	<b>then</b> Humidity = Normal <b>and</b> Play = Yes	4/8
<b>If</b> Play = Yes	<b>then</b> Humidity = Normal <b>and</b> Windy = False	4/9
		4/14
<b>If</b> True	<b>then</b> Humidity = Normal <b>and</b> Windy = False <b>and</b> Play = Yes	

Rules with same itemset: all  $k-1$  consequent rules have at least the same confidence, as  $k$  consequent rules

# Generating rules efficiently

- We are looking for all high-confidence rules
  - Which items on the RHS of the rule?
  - But: brute-force method is  $(2^N - 1)$
- Observation:  $(c + 1)$ -consequent rule can only have a confidence above the threshold if all corresponding  $c$ -consequent rules also have
- Building  $(c + 1)$ -consequent rules from  $c$ -consequent ones!

# Example

- 1-consequent rules:

**If** Outlook = Sunny **and** Windy = False **and** Play = No  
**then** Humidity = High (2/2)

**If** Humidity = High **and** Windy = False **and** Play = No  
**then** Outlook = Sunny (2/2)

- Corresponding candidate 2-consequent rule:

**If** Windy = False **and** Play = No  
**then** Outlook = Sunny **and** Humidity = High (2/2)

- Final check of antecedent against data!

# Discussion of association rules

- Above method makes one pass through the data for each different size item set
  - Other possibility: generate  $(k+2)$ -item sets just after  $(k+1)$ -item sets have been generated
    - Result: more  $(k+2)$ -item sets than necessary will be considered but less passes through the data
- Practical issue: generating a certain number of rules (e.g. by incrementally reducing min. support)



# Summary/concepts/questions

- 1R
  - Model
  - Algorithm
  - Unknown values
  - Numeric attributes
- Naive Bayes
  - Model
  - Algorithm (Learning, predicting)
  - Laplace estimator
  - Unknown values
  - Numeric attributes
- ID3 decision tree
  - Level of surprise, entropy, information gain, algorithm
- Rules: problem with interpretation

# Summary/concepts/questions - 2

- Linear models
  - Form of linear model for regression
  - Form of error to be minimized
- Linear models for classification
  - Logistic regression
  - Perceptron
  - Winnow
  - Instance based (Nearest neighbour) learning
    - „Knowledge representation”
    - Algorithm
    - Problems
- Clustering
  - Knowledge representation
  - K-means: algorithm + limitations
- Association rule mining
  - Knowledge representation
  - Item, itemset, support, confidence, lift
  - Apriori algorithm