



DM & ML

Ensemble learning

Gergely Lukács

Pázmány Péter Catholic University

Faculty of Information Technology

Budapest, Hungary

lukacs@itk.ppke.hu

Combining multiple models

„Crowd can be smarter than the individuals”

- Basic idea of “meta” learning schemes: build different “experts” – or crowd -- and let them vote
 - Sufficiently diverse crowd...
- Advantage:
 - often very good predictive performance
- Disadvantage:
 - interpretation of model ? (black box...)

Stacking / 1

- Train different machine learning models on same dataset
 - *Level-0 Models* (Base-Models)
 - (Different algorithms)
 - Standard ML algorithms (e.g., Decision tree, Bayesian network)
 - Special algorithms of the application area (e.g. Natural Language Processing, parsing)
 - Models of different teams in competitions („team merging”)
- Combine the results of those models
 - *Level-1 Model* (Meta-Model)
 - Often simple model (in principle: any model)
 - David Wolpert: “relatively global, smooth” model
 - Base learners do most of the work
 - Reduces risk of overfitting

Stacking / 2


- Output of Level-0 models
 - Classification: label
 - Even better: also probabilities!
 - Numeric prediction („Regression”): Value
 - These are used to train Level-1 model
 - Optional: additionally attributes of data as additional context

Stacking / 3

- Data used for training level-0 models not to be used to train level-1 model! – overfitting!
 - Cross-validation-like scheme is employed
- Training of Level-0 models
 - Some sort of validation, typical: cross validation
 - for training Level-1 model
 - Finally: all data used for training
- Training of Level-1 model
 - Data not used to train Level-0 models

Bagging

- **Same algorithm** (learning scheme) for all classifiers
- **Different datasets**
- “Idealized” version:
 - Sample several training sets of size n (instead of just having one training set of size n)
 - Build a classifier for each training set
- Combining predictions by voting/averaging
 - Simplest way (e.g. each model receives equal weight)
 - Reduces variance by voting/averaging, thus reducing the overall expected error



Bagging – creating several datasets

- Problem: we only have one dataset!
- Solution: generate new datasets of size n by **sampling with replacement** from original dataset
 - B-Agging = Bootstrap aggregating

Bagging classifiers




model generation

```
Let n be the number of instances in the training data.  
For each of t iterations:  
    Sample n instances with replacement from training set.  
    Apply the learning algorithm to the sample.  
    Store the resulting model.
```

classification

```
For each of the t models:  
    Predict class of instance using model.  
Return class that has been predicted most often.
```


Bagging



Same algorithm! (e.g. decision tree, Bayesian Network,.... anything)

Combining results:
voting / averaging

More on bagging

- Learning scheme is unstable
(Small change in training data can make big change in model; e.g. decision trees)
-> almost always improves performance
- Theoretically, bagging would reduce variance without changing bias
 - In practice, bagging can reduce both bias and variance
 - For high-bias classifiers, it can reduce bias
 - For high-variance classifiers, it can reduce variance
 - Usually, the more classifiers the better
 - Can help a lot if data is noisy
 - In the case of classification there are pathological situations where the overall error might increase

Bagging with costs

- good probability estimates
 - ♦ instead of voting, the individual classifiers' probability estimates are averaged
- also for learning problems with costs

Randomization

- Can randomize learning algorithm instead of data
- Some algorithms already have a random component:
 - e.g. initial weights in neural net
- Most algorithms can be randomized, e.g. greedy algorithms:
 - ♦ Pick from the N best options at random instead of always picking the best options
 - ♦ E.g.:
 - ♦ attribute selection in decision trees
 - ♦ *random subspaces*: random subsets of attributes in nearest-neighbor scheme
- Can be combined with bagging

Random forest

- „Forest” – multiple decision trees
- Random:
 - Feature randomness:
Only considers attributes in a random subset of all attributes for a split (instead of all attributes)
 - Data randomness:
usually bootstrapping data (sampling with replacement)
- Voting

Rotation forest

accurate and diverse ensemble members

(2006, Rodríguez et al.)

- Bagging creates ensembles of accurate classifiers with relatively low diversity
 - ♦ Bootstrap sampling creates training sets with a distribution that resembles the original data
- Randomness in the learning algorithm increases diversity but sacrifices accuracy of individual ensemble members
- Accuracy-diversity dilemma
- Rotation forests have the goal of creating accurate **and** diverse ensemble members

Rotation forest/2

- An iteration involves
 - ♦ Randomly dividing the input attributes into k disjoint subsets
 - ♦ Applying PCA to each of the k subsets in turn
 - ♦ (all attributes are kept; **it preserves all information!**)
 - ♦ Learning a decision tree from the k sets of PCA directions

Rotation forest/3

- ♦ „Rotation forest”
 - ♦ Base classifier: decision tree → „forest”
 - ♦ PCA is a simple rotation of the coordinate axes → „rotation”
- ♦ Rotation forest has the potential to improve on diversity significantly without compromising the individual accuracy


Boosting

- Boosting: explicitly seeks models that complement one another
- Works well with weak models


Boosting 2

- Similar to Bagging
 - Voting (classification), averaging (numeric prediction)
 - Models of the same type (e.g. decision trees)
- Different from bagging
 - Iterative, models not built separately
 - New model is encouraged to become expert for instances classified incorrectly by earlier models
 - Intuitive justification: models should be experts that complement each other
 - Weighted voting!
- There are several variants of this algorithm


Boosting illustration



Boosting illustration




Boosting illustration




Boosting illustration

*Weights
Increased*




Boosting illustration



Boosting illustration

*Final classifier is
a combination of
single classifiers*



AdaBoost.M1

Model Generation

Assign equal weight to each training instance.

For each of t iterations:

- Apply learning algorithm to weighted dataset and store resulting model.

- Compute error e of model on dataset and store error.

- If e equal to zero, or e greater or equal to 0.5

 - Terminate model generation.

- For each instance in dataset:

 - If instance classified correctly by model

 - Multiply weight of instance** by $e / (1 - e)$.

- Normalize weight for all instances.

Classification

Assign weight of zero to all classes.

For each of the t (or less) models:

- Add $-\log(e/(1-e))$ to **weight of class** model predicts.

Return class with highest weight.

More on boosting

- Theoretical result:
 - Upper limit for training error decreases exponentially
 - Boosting works with weak learners only condition: error doesn't exceed 0.5
 - Weak learner: only slightly better, than random predictor
- Practical issue: Boosting needs weights
 - adapting learning algorithm to use weights
 - resample with probability determined by weights
- AdaBoost in fact
 - Additive model
 - Particular loss function (exponential loss)

(17. Learning: Boosting)

- <https://www.youtube.com/watch?v=UHBmv7qCey4&t=2752s>

Gradient Boosting

- AdaBoost: adjusting weights of data points...
- Difference between the prediction and the ground truth.
 - Optimization problem
 - on suitable cost function
 - over function space
 - Iterative, greedy
 - by choosing function pointing in negative gradient direction

Elements of Learning

Example:

Linear regression

$$\hat{y}_i = \sum_j \theta_j x_{ij}$$

- Model
 - Parameters
- Objective Function
 - Training loss
 - + regularisation!!!

$$L(\theta) = \sum_i (y_i - \hat{y}_i)^2$$

$$\text{obj}(\theta) = L(\theta) + \Omega(\theta)$$

XGBoost – theoretical advantages

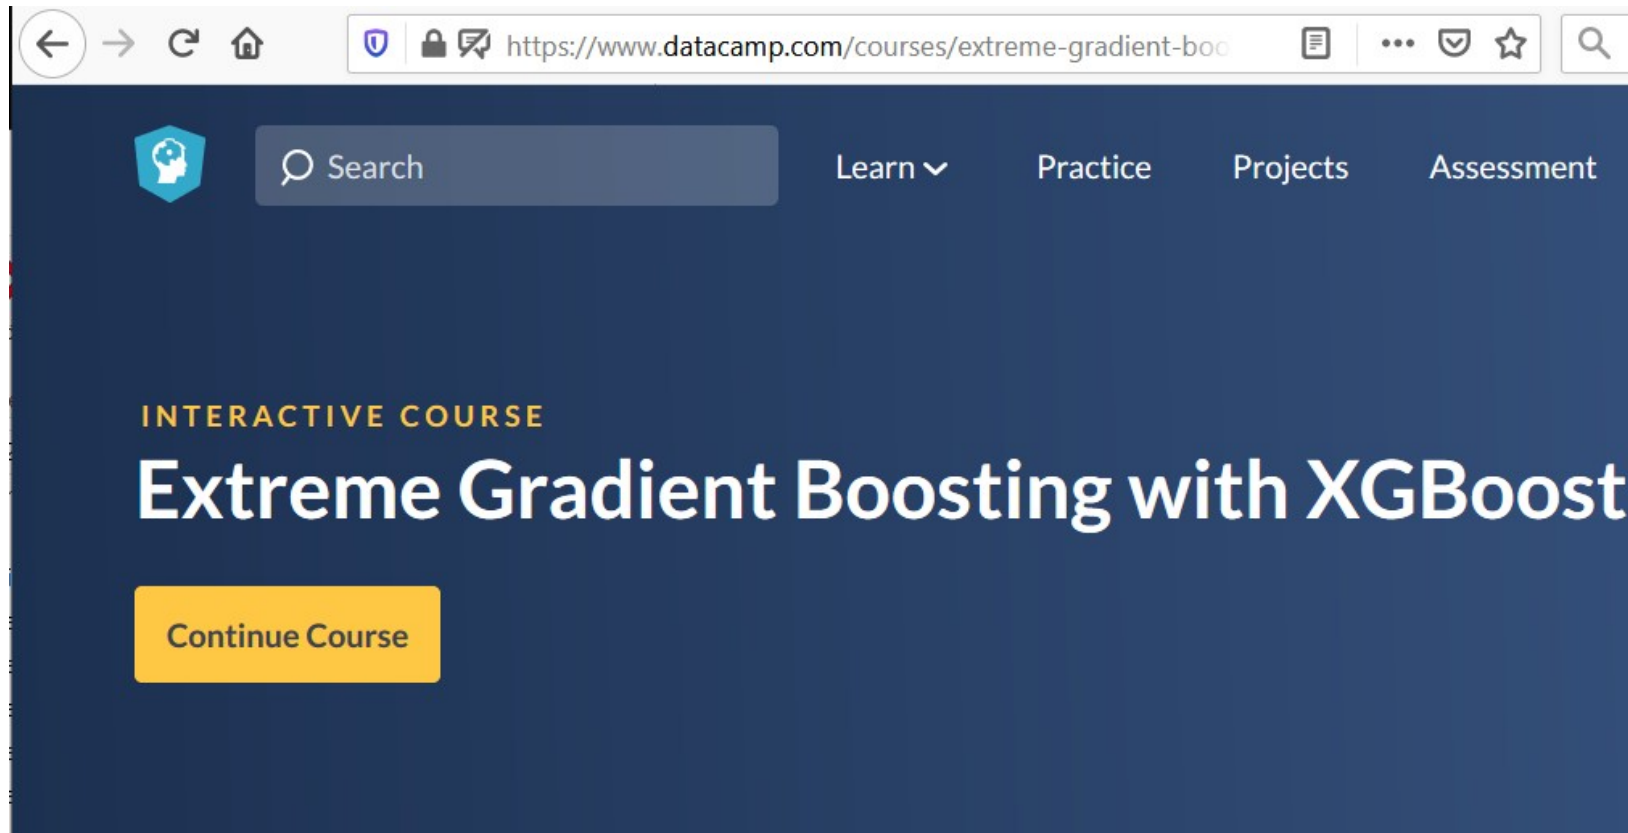
- Decision tree ensembles
- Regularisation!
 - Formalizes complexity of tree classifiers
 - overfitting reduced!!
- Custom optimization objectives and evaluation criteria
- +... (handling missing values, built-in cross validation, continue existing model...)

XGBoost – computing advantages

- Use of **sparse matrices** with sparsity aware algorithms
- Improved **data structures** for better **processor cache utilization** which makes it faster.
- Better support for **multicore processing** which reduces overall training time.


XGBoost – Parameter Tuning!

- <https://www.analyticsvidhya.com/blog/2016/03/complete-guide-parameter-tuning-xgboost-with-codes-python/>



The image is a screenshot of a web browser displaying the DataCamp website. The browser's address bar shows the URL `https://www.datacamp.com/courses/extreme-gradient-boo`. The website's header is dark blue and contains the DataCamp logo (a teal shield with a white brain icon) on the left. To the right of the logo is a search bar with the placeholder text "Search". Further right are navigation links: "Learn" with a dropdown arrow, "Practice", "Projects", and "Assessment". The main content area has a dark blue background. It features the text "INTERACTIVE COURSE" in a small, yellow, all-caps font. Below this is the course title "Extreme Gradient Boosting with XGBoost" in a large, white, sans-serif font. At the bottom left of this section is a yellow rectangular button with the text "Continue Course" in a dark font.

← → ↻ 🏠 🔒 `https://www.datacamp.com/courses/extreme-gradient-boo` 📄 ⋮ 🛡️ ☆ 🔍

 Search Learn ▾ Practice Projects Assessment

INTERACTIVE COURSE

Extreme Gradient Boosting with XGBoost

Continue Course

XGBoost vs Deep learning

- Start with simple models!!
- XGBoost
 - Easier to train
 - Less computational resources
 - Better if categorical + numeric features
- Deep learning
 - Image Recognition, Computer Vision, Natural Language Processing (some sort of structure, space)

Comparison

Stacking	Bagging	Boosting
Data1 = Data2 = ... = Data m	Data1 \neq Data2 \neq ... \neq Data m	
	Resample training data	Reweight training data
Learner1 \neq Learner2 \neq ... \neq Learner m	Learner1 = Learner2 = ... = Learner m	
Level-1 model	Vote	Weighted vote

Summary/concepts/questions

- Bias-variance decomposition
- Stacking
- Bagging
- Randomization
 - Random Forest
 - Rotation forest
- Boosting
 - AdaBoost
 - Gradient Boosting
 - XGBoost